

COMP3018: Report (Literature Review & Programming Project)

Cognitive Robotics

Contents

TODO	1
VERIFY PAGE NUMBERS (check each against the actual PDF)	2
Mentor	8
DR ALY QUESTIONS:	8
First-to-do:	8
General:	8
1- Task (3) Literature Review	10
[] 1.3.1 Therapeutic and Emotional Support	11
[] 1.3.2 Medication Adherence and Daily Living Support	12
[] 1.3.3 Physical Rehabilitation and Mobility	12
[] 1.4.1 Challenges	13
[] 1.4.2 Ethical Implications	13
1.6. Task-3 References	15
1.6.1 Papers and Journal Articles	15
2- Task (4) Novel Programming Project	16
2.1. Introduction (10%; Salman)	16
2.2. Background (10%; Alfie)	17
2.3. Methods & Setup (35%; Alfie)	17
2.3.1 System Architecture	17
2.3.2 Input Layer: Multimodal and Behavioural Signals	17
2.3.3 Process Layer: Multi-Signal State Inference	21
2.3.4 Generate Layer: Dynamic Prompt Construction	24
2.3.5 Output Layer: Aligned Multimodal Response	24
2.3.6 Adaptation Self-Evaluation and Session Persistence	24
2.4. Outcome & System Analysis (30%; Salman)	24
2.5. Conclusion (10%; Alfie)	24
2.6 Task-4 References (5%)	24
Alfie's	25
2.7. Appendix	26
2.7.1. Code	26
2.7.2. Video Demo	26
Appendices	27
Appendix A: AI Declaration	27

TODO

- ☐ verify word count

VERIFY PAGE NUMBERS (check each against the actual PDF)

Vernon, Metta and Sandini (2007) – papers/Vernon, Metta and Sandini (2007) - A Survey of Artificial Cognitive Systems.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	298, 304	S1.1, S1.2	Vernon et al. (2007, p. TODO)	Try p. 155	Section “What is Cognition?”; cognition cycle diagram (anticipate, learn, adapt + perception/action)
- []	306	S1.2	Vernon et al. (2007, p. TODO)	Try p. 163	Memory section; episodic vs semantic memory distinction (Tulving’s taxonomy)

Sciutti et al. (2023) – papers/Sciutti et al. (2023) - The Present and the Future of Cognitive Robotics.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	298, 304, 314, 509	S1.1- S2.2	Sciutti et al. (2023, p. 160)	p. 160	“flexible, context-sensitive action, knowing what they are doing and why”; “reason about their actions and modify their behavior”
- []	302	S1.2	Sciutti et al. (2023, p. 160) – “intersection of Robotics, AI, Cognitive Sciences”	p. 160	CHECK: this exact phrase may NOT be here; may come from Sandini, Sciutti & Vernon (2021) encyclopaedia entry instead
- []	350	S1.5	Sciutti et al. (2023, pp. 162-163)	pp. 162-163	CHECK: “integrating machine learning techniques with model-based approaches” – may only be on p. 162, not spanning 163

Tapus, Matarić and Scassellati (2007) – papers/Tapus, Matarić and Scassellati (2007) - Socially Assistive Robotics.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	312	S1.3.1	Tapus et al. (2007, p. TODO)	Try p. 35	PARO listed: “robotic animal toys, such as a seal (i.e., PARO [2])”
- []	509	S2.2	Tapus et al. (2007, p. 35)	p. 35	“helping human users through social rather than physical interaction”

Wada and Shibata (2007) – papers/Wada and Shibata (2007) - Living With Seal Robots.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	312, 342	S1.3.1, S1.4.2	Wada and Shibata (2007, p. 974)	p. 974	NEEDS MANUAL CHECK – download PDF; look for PARO reducing agitation / mood improvement in elderly

Fong, Nourbakhsh and Dautenhahn (2003) – papers/Fong, Nourbakhsh and Dautenhahn (2003) – A Survey of Socially Interactive Robots.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	314	S1.3.1	Fong et al. (2003, p. 145)	p. 145	Section 1.2; Breazeal’s four classes of social robots; “shallow models of social cognition” under Social Interface
- []	326	S1.3.3	Fong et al. (2003, p. 149)	p. 149	Section 2.3 Embodiment; “mutual perturbation” / “perturbatory channels”
- []	509, 538	S2.2, S2.3.3	Fong et al. (2003, p. 148)	p. 148	CHECK: p. 148 covers design issues, NOT emotion recognition. Try p. 155 (human-oriented perception listing) or p. 156 (speech/facial emotion analysis)

Lee and See (2004) – papers/Lee and See (2004) – Trust in Automation Designing for Appropriate Reliance.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	318	S1.3.2	Lee and See (2004, p. 54)	p. 54	Definition in italics: “the attitude that an agent will help achieve an individual’s goals in a situation characterized by uncertainty and vulnerability”

Hancock et al. (2011) – papers/Hancock et al. (2011) – A Meta-Analysis of Factors Affecting Trust in Human-Robot Interaction.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	320, 334	S1.3.2, S1.4.1	Hancock et al. (2011, p. 522) – performance strongest predictor	p. 522	Table 1; “performance factors were more strongly associated ($r = +0.34$)”; Cohen’s $d = +0.71$
- []	318	S1.3.2	Hancock et al. (2011, p. 522) – “observation cannot reliably disambiguate”	p. 522	CHECK: this POMDP-style claim may NOT appear anywhere in Hancock et al.; possibly misattributed
- []	334	S1.4.1	Hancock et al. (2011, p. 522) – “29 studies, modest variance”	pp. 520-522	29 studies stated on p. 520; overall $r = +0.26$ on p. 521; paper uses “moderate” not “modest”

Chen et al. (2020) – papers/Chen et al. (2020) – Trust-Aware Decision Making for Human-Robot Collaboration.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	320	S1.3.2	Chen et al. (2020, p. 6)	p. 6 (article page :6)	Section 3.4 “Maximizing team performance”; Fig. 3 (Trust-POMDP model); “We maintain a belief b over the human’s trust”

Garcez and Lamb (2023) – papers/Garcez and Lamb (2023) – Neurosymbolic AI The 3rd Wave.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- [	320, 511, 663	S1.3.2, S2.2, S2.5	Garcez and Lamb (2023, p. 12389)	Local PDF is arXiv preprint (pp. 1-28); p. 12389 = journal page 3 of published version	CHECK against published <i>AI Review</i> version (journal pp. 12387-12406). On ~p. 3: third wave / neural-symbolic labour division

Nikolaidis, Hsu and Srinivasa (2017) – papers/Nikolaidis, Hsu and Srinivasa (2017) – Human-Robot Mutual Adaptation in Collaborative Tasks.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- [	320, 624	S1.3.2, S2.3.3	Nikolaidis et al. (2017, p. 625)	p. 625	“69 samples”; “U = 180, p = 0.048”; MOMDP mutual-adaptation condition
- [	334	S1.4.1	Nikolaidis et al. (2017, p. 627)	p. 627	MOMDP discussion; r = -0.066 (no correlation between trustworthiness and inferred adaptability)

Brooks (1991) – papers/Brooks (1991).pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- [	326	S1.3.3	Brooks (1991, p. TODO)	??	NEEDS MANUAL CHECK – find argument that intelligence emerges from physical interaction, not abstract representation

Mataric et al. (2007) – papers/Mataric et al. (2007).pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- [	326	S1.3.3	Mataric et al. (2007, p. TODO)	??	NEEDS MANUAL CHECK – find stroke survivors engaging more with embodied robot than screen-based alternative

Tapus, Tăpuş and Mataric (2008) – papers/Tapus, Tapus and Mataric (2008).pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- [	326	S1.3.3	Tapus, Tăpuş and Mataric (2008, p. TODO)	??	NEEDS MANUAL CHECK – find adaptive personality matching (interaction distance/speed) improving task performance

Papadimitriou and Tsitsiklis (1987) – papers/Papadimitriou and Tsitsiklis (1987) - The Complexity of Markov Decision Processes.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	332	S1.4.1	Papadimitriou and Tsitsiklis (1987, p. 448)	p. 448	Theorem 6: “The partially observed problem is PSPACE-hard”

Pineau, Gordon and Thrun (2003) – papers/Pineau, Gordon and Thrun (2003) - Point-Based Value Iteration An Anytime Algorithm for POMDPs.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	332	S1.4.1	Pineau et al. (2003, p. 1025)	p. 1025	First page / abstract: “Point-Based Value Iteration (PBVI) algorithm”

Kaelbling, Littman and Cassandra (1998) – papers/Kaelbling, Littman and Cassandra (1998) - Planning and Acting in Partially Observable Stochastic Domains.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	332, 663	S1.4.1, S2.5	Kaelbling et al. (1998, p. 120)	p. 120	CHECK: p. 120 is the tiger problem (Section 5.1, toy example). For POMDP framework definition try p. 105; for belief-state planning try p. 108

Silver and Veness (2010) – papers/Silver and Veness (2010) - Monte-Carlo Planning in Large POMDPs.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	332	S1.4.1	Silver and Veness (2010, p. 1)	p. 1	Abstract: “Monte-Carlo algorithm for online planning in large POMDPs”

Broadbent, Stafford and MacDonald (2009) – papers/Broadbent et al. (2009).pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	334	S1.4.1	Broadbent et al. (2009, p. TODO)	??	

Desai et al. (2013) – papers/Desai et al. (2013) - Impact of Robot Failures and Feedback on Real-Time Trust.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	334, 538	S1.4.1, S2.3.3	Desai et al. (2013, p. 256)	p. 256	CHECK: local PDF is conference format with no printed page numbers. Verify against HRI 2013 proceedings. Look for: trust drops after reliability failures, slow recovery

Wachter, Mittelstadt and Floridi (2017) – papers/Wachter, Mittelstadt and Floridi (2017) – Why a Right to Explanation Does Not Exist in the GDPR.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	340	S1.4.2	Wachter et al. (2017, p. TODO)	Try p. 76 (abstract) or p. 82 (Section 3 heading)	p. 76: “GDPR does not implement a right to explanation”; p. 82: Section heading “Why there is no ‘right to explanation’ ”

Sharkey (2014) – papers/Sharkey (2014) – Robots and Human Dignity.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	342	S1.4.2	Sharkey (2014, p. 6 (VERIFY))	Manuscript p. 6 = journal p. 68	“a robot that dealt impersonally with an older person, without knowing or using their name or their preferences...” in Nordenfelt Dignity of Identity context

Sharkey and Sharkey (2012) – papers/Sharkey and Sharkey (2012) – Granny and the Robots Ethical Issues in Robot Care for the Elderly.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	350	S1.5	Sharkey and Sharkey (2012, p. 27)	p. 27	Abstract/introduction; concerns about reducing human contact. Detailed argument is on pp. 30-31

Ahn et al. (2022) – papers/Ahn et al. (2022) – Do As I Can Not As I Say Grounding Language in Robotic Affordances.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	511	S2.2	Ahn et al. (2022, p. 1)	p. 1	Abstract: “constrain the model to propose natural language actions that are both feasible and contextually appropriate”

Smedegaard (2019) – papers/Smedegaard (2019) – Reframing the Role of Novelty within Social HRI from Noise to Information.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	511, 663	S2.2, S2.5	Smedegaard (2019, p. 4)	p. 4 (= proceedings p. 414)	Novelty as “original feature of experience”. CHECK: the novelty-fading engagement claim may be stronger on p. 2 (proceedings p. 412)

Ji et al. (2023) – papers/Ji et al. (2023) - Survey of Hallucination in Natural Language Generation.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- []	647	S2.3.4	Ji et al. (2023, p. 3)	p. 3	“deep learning based generation is prone to hallucinate unintended text”
- [X]	596	S2.3.2	Ji et al. (2023, p. 3)	p. 3	<i>degeneration</i> defined; output that is “bland, incoherent, or gets stuck in repetitive loops” (§1 INTRODUCTION, para. 2, sent. 3). verified 2026-04-22 via PDF extraction (cross-checked with Gemini)

Picard (1997) – papers/Picard (1997) - Affective Computing.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- [X]	517	S2.2	Picard (1997, p. 1, Abstract)	p. 1	Abstract: “affective computing,” computing that relates to, arises from, or influences emotions” (full founding definition). verified 2026-04-14 via PDF extraction

Spezialetti, Placidi and Rossi (2020) – papers/Spezialetti, Placidi and Rossi (2020) - Emotion Recognition for Human-Robot Interaction.pdf

OK?	Line(s)	Section	Citation as written	Go to page...	You should see...
- [X]	517	S2.2	Spezialetti, Placidi and Rossi (2020, pp. 1-2, Introduction)	pp. 1-2	Introduction; bullet “Ability of robots to infer the human emotional state” (p. 2); reviews recognition across facial, vocal, brain, peripheral physiological channels. verified 2026-04-14

Radford et al. (2023) – papers/Radford et al. (2023) - Robust Speech Recognition via Large-Scale Weak Supervision.pdf

OK?	Line(s)	Citation as written	Go to page...	You should see...
- [X]	544	S2.3.2 Radford et al., 2023, Abstract)	p. 1	Abstract (final sentence): “When compared to humans, the models approach their accuracy and robustness”. No inline page cited (PMLR volume-pagination not printed on page). verified 2026-04-14

-
- ☐ if enough time: emergent cognitive robot architecture
 - ☐ cite Iuliia Kotseruba · John K. Tsotsos
 - ☐ use ‘inference’ naturally, i.e., not so perfect that it is likely ai generated but instead slightly not 100% correct like a human would do, ygm?

Mentor

Scenario 1) patient takes meds = + 100 points Scenario 2) Patient is annoyed = -10 points When you implement this you need to be careful to structure the rewards such that the system does not pester the patient as it seeks to maximise its rewards You’ll need to play around with the precise ratios between 1 (behaviour you want to encourage) and 2 (behaviour you want to discourage) to achieve the desired attitude from the robot

DR ALLY QUESTIONS:

- ☐ so maths and code won’t affect word count right? However I never explicitly knew if captions within a figure-diagram count towards word count?
-

First-to-do:

- ☐ **Most critical:** verify all page numbers and sentences manually
- ☐ **Most critical:** verify all links to papers manually
- ☐ do all page number TODOs
- ☐ implement loads of peer-reviewed papers everywhere again

General:

- ☐ ensure no overused complex words
- ☐ Download Broadbent via Plymouth library
- ☐ decides the reward based on what it observes = inferring the reward
- ☐ discuss POMDP maths
- ☐ use wording from 3018 Task-4 Proposal Google Doc
- ☐ talk about the LfD (learning-from-demonstration)
- ☐ talk about IRL (inverse reinforcement learning)
- ☐ relatively talk about how it relates to others, motivation
- ☐ CRAMS figure verified – 5/6 actions appear (Back_Off absent because true state never sustained Low Trust long enough). Update report Task 4 discussion to: 1) explain why the action timeline shows context-sensitive selection (link each action cluster to the reward territory that produced it), 2) note

Back_Off correctly absent given the Medium-trust initial state and stress profile, 3) highlight META-ADAPT triggers (red dotted lines) as evidence of metacognition detecting the stress event within 2 steps

- ☒ USE ‘misclassified’- [] consider a project wherein it is ‘cognitive robotics’ (lecture 9) ensure it involves what we have learnt in the labs ‘
- ☐ ‘persons’
- ☐ make the robot kinda like how I disussed you should make it in the set exercises
- ☐ discuss mathematical notation for the POMDP stuff??? (if not done in set exercises)
 - ☐ maths and diagrams affect wordcount?
 - ☐ Trust-POMDP diagram
 - ☒ Cite: Chen, M., Nikolaidis, S., Soh, H., et al., “Trust-Aware Decision Making for Human-Robot Collaboration: Model Learning and Planning”, ACM Trans. Hum.-Robot Interact., 9(2), 2020
 - ☐ model trust in lates diabram? or just latex the fundamental diagtam of POMDP similar to lecture slied in POMDP lectures (10,11)
 - ☐ latex diagram of continous state in POMDP similar to non-monotnoric graph (not about non-monitonic itsekf just a similar graph)
 - ☐ utilise lec-7 for POMDP insights; similar to machine learning set exercises where i give a quick breakdown if word count allowance
- ☐ args and kwargs (if i can do this)
- ☒ peer-reviewed or conference papers
- ☐ In this section, you should focus on providing enough description of the supervised learning, neural network, and naïve Bayes models.
- ☐ Do not assume the reader knows the basics. Dedicate specific paragraphs to explicitly defining the algorithms and the broader category (Supervised Learning) before diving into your implementation.
- ☐ Then, refer to some studies that have utilised neural networks and naïve Bayes models in your area using the selected database
- ☐ Ensure your literature review in the introduction explicitly cites papers that use your specific dataset (or very similar ones), establishing a clear baseline before you begin
- ☐ TODO: talk about the LfD (learn-from-demonstration)
 - ☐ TODO: talk about IRL (inverse reinforcement learning)
- ☐ TODO: talk about the LfD (learn-from-demonstration)
 - ☐ TODO: talk about IRL (inverse reinforcement learning)

LECTURE 9 – TODOs FOR ASSESSMENT 2 REPORT Task 3: Literature Review (Assistive Robotics Essay): - [X] Frame assistive robotics as requiring cognitive capabilities, not just social behaviour – use Aly’s hierarchy: “intelligence deployed over the social layer, not vice versa” - [X] Reference Sciutti et al. (2023) definition of cognitive robotics (replaced Cangelosi book – not peer-reviewed) - [X] Use Vernon, Metta and Sandini (2007) cognition cycle (replaced Vernon 2014 book – not peer-reviewed) - [X] Discuss theory of mind, prospection, and episodic memory as future directions/current gaps in assistive robotics - [X] Acknowledge the 42-definitions problem to argue the field still lacks consensus on what cognition actually is (shows critical thinking) - [X] Connect embodied cognition (“intelligence means body”) to ethical implications of robots entering intimate care spaces Task 4: Programming Project - [X] GET APPROVAL FROM ALY – confirm extending set exercises POMDP into cognitive architecture implementation is acceptable; confirm simulation-only is fine - [X] Wait for Lectures 10/11 on cognitive architectures before finalising design – Aly said the building blocks reappear in those lectures - [] Add metacognition module (new from this lecture) – system monitors its own reasoning, flags when repeated actions produce negative outcomes - [] Implement explicit episodic vs semantic memory distinction – Aly specifically said “please distinguish or remember these two”

- [] Label trust inference as theory of mind explicitly - [] Label POMDP planning over future states as prospection explicitly - [] Label observation filtering as attention (selective/suppressive) - [] Frame all actions as goal-directed (Aly’s term for purposeful cognitive actions vs reactive behaviour) - [] Frame project as cognitive robotics (not just social robotics) – Aly said nobody has ever done this; “invitation for challenging minds” - [] Don’t skip definitional rigour – define cognitive robotics, cognition, and key terms precisely in the Background section - [] 5-min video: walk through a scenario showing perceive -> attend -> reason -> act -> learn -> adapt cycle in action

1- Task (3) Literature Review

☐ cite this lecture and book: /Users/alfienurse/Desktop/gitdev/Uni/Undergrad/year-3/Work/1st Year/Semester 2/3018-cw/learning/[] final: cognitive architecture

☒ [X] ## 1.1. Introduction

Ageing populations and a shrinking care workforce positioned **assistive robotics** (*human-supportive robots within physical, cognitive, and social realms of impairment affecting daily-living activities*); a prominent technological response to a widening care gap. The field spans physical prosthetics; surgical assistance; neurodivergent support; exoskeletal rehabilitation, and social companionship; this essay however focuses on *socially* assistive robotics (*SAR*), an assistive-robot subfield wherein the robot “focuses on helping human users through social rather than physical interaction” (Tapus, Matarić and Scassellati, 2007, Abstract), as here most active research and ethical tension converge.

Robots now administer medication reminders, facilitate rehabilitation exercises, and therapeutically companionate in clinical and domestic settings: reduced caregiver burden, improved patient outcomes wherein residents became “more active and more communicative, both with each other and their caregivers” (Wada and Shibata, 2007, p. 973), and increased “social interaction” among elderly residents (Wada and Shibata, 2007, p. 972, Abstract).

However, most-current assistive robots operate at what Sciutti et al. (2023, p. 160) call the social layer: they react to immediate stimuli but lack the cognitive depth to anticipate user needs, remember past interactions, or reason about their own performance. Sciutti et al. argue that effective assistive robots should be *cognitive*: capable of “flexible, context-sensitive action, knowing what they are doing and why they are doing it.” Vernon, Metta and Sandini (2007, p. 151) formalise this requirement via a “virtuous cycle that is embedded in an ongoing process of action and perception” (the agent anticipates → learns → adapts to achieve autonomy). This essay contends that assistive robotics should graduate from reactive social behaviour to cognitive capability (intelligence deployed *over* the social layer) if it is to deliver sustained, personalised support. The following sections survey the theoretical foundations thereof: evaluate prominent applications through this cognitive lens, discuss challenges and ethical implications, and identify future directions.

[X] ## 1.2. Literature Review

Cognitive robotics: defined by Sciutti et al. (2023, p. 160), lies at the intersection of Robotics, Artificial Intelligence, and Cognitive and Biological Sciences, combining “sensorimotor behaviour, higher-level functions, and social capabilities of an intelligent robot.” This interdisciplinary grounding distinguishes it from conventional robotics (*treats the robot as purely engineered*) and from social robotics (*addresses interaction behaviour without necessarily modelling cognitive processes*). The distinction is consequential: a robot that smiles when a patient smiles is social; a robot that infers *why* the patient is smiling, and adjusts its future strategy accordingly, is therefore *cognitive*.

☐ verify following lecture slides

Vernon, Metta and Sandini (2007, p. TODO) synthesise the field’s definitional plurality into a core cycle. The European Network for Advancement of Artificial Cognitive Systems (euCognition) catalogued 42 definitions of cognition, yet the common thread therein is: anticipation, learning, and adaptation, intersected with perception and action to create autonomy (Fig. 1). This cycle provides an architectural checklist for assistive robots: a system that cannot direct its gaze toward relevant stimuli whilst suppressing irrelevant ones

(*selective attention*), anticipate the outcome of its actions (*i.e. prospection*), learn from past interactions (*memory*), or adapt its strategy when performance declines (*metacognition*) is, per this framework, not yet cognitive. Sciutti et al. (2023, TODO: verify p. 160) further specify that cognitive robots should “reason about their actions and modify their behavior to improve their effectiveness”; a capacity termed *theory of mind*, wherein the agent infers another’s hidden mental state from observable behaviour.

Furthermore, memory is not monolithic. Vernon, Metta and Sandini (2007, p. TODO) distinguish *episodic memory* (*records of specific past experiences and their contextual outcomes*) from *semantic memory* (*general knowledge about the world, including spatial relationships and factual constraints*); a typology Laird (2012, p. 225) crystallises in implementational terms: episodic memory is “what you ‘remember’” whilst semantic memory is “what you ‘know’”. For example, assistive-medication robots need episodic memory to recall that a user refused medication after a restless night, and semantic memory to know certain drugs cannot also be administered. Whilst the 42-definitions problem confirms the field lacks consensus on what cognition *per se* is, the common thread (anticipation, learning, adaptation) is exactly what assistive robotics demands.




Figure 1: Soar’s processing cycle, adapted from Laird (2012, fig. 4.7, p. 79), as a concrete instantiation of Vernon, Metta and Sandini’s (2007) cognition cycle. Laird specifies that the cycle “consists of four phases: Input, Operator Selection, Operator Application, and Output” (Laird, 2012, p. 79), wherein Input maps to perception, Operator Selection draws on episodic and semantic memory to anticipate consequences (*i.e. prospection*), Operator Application executes the chosen behaviour through parallel rule-firing waves, and Output returns the agent to the environment whilst the loop hands control back to Input for continual adaptation. This architecture exposes the deficit of reactive systems: PARO, for instance, implements only Input → Output, lacking the substate-driven Operator Selection wherein deliberation over latent user states could occur. Furthermore, Laird argues that “the most significant effect of chunking is that it eliminates processing in substates for situations similar to ones experienced in the past” (Laird, 2012, p. 164), wherein repeated deliberations compile into procedural rules; a mechanism distinct from POMDP belief-space approximation, yet complementary thereto in addressing real-time embodied operation under care-time constraints.

□ ## 1.3. Applications

[] 1.3.1 Therapeutic and Emotional Support

The PARO therapeutic seal robot represents one of the most-widely deployed platforms within socially assistive robotics (Tapus, Matarić and Scassellati, 2007, p. TODO). Wada and Shibata (2007, p. 974) demonstrate that PARO reduces agitation and improves mood in patients with dementia, utilising tactile sensors and auditory processing to modulate its behaviour in response to touch and voice. Clinical trials report that urinary stress indicators “significantly improved” after PARO’s introduction (Wada and Shibata, 2007, TODO: verify: p. 978, Table II), therefore the platform has been adopted in care homes across Japan, Europe, and the United States.

Notwithstanding the benefits, PARO operates at the reactive layer: it possesses no theory of mind (it cannot infer *why* a patient is agitated (loneliness, pain, confusion) nor episodic memory of *what* calmed the patient previously. A cognitively-equipped therapeutic robot, in contrast, would anticipate mood shifts via prospection → recall that music soothed this patient yesterday via episodic memory → adapt its strategy via metacognition. Insofar as PARO’s effectiveness plateaus because it cannot personalise its responses over time, the cognitive gap is not just theoretical but potentially clinically consequential. The architectural cost is precise: without episodic memory, the agent’s perception is “limited both spatially and temporally—that

is, to the here and now” (Laird, 2012, p. 233), wherein each session resets to the immediate present. Fong, Nourbakhsh and Dautenhahn (2003, p. TODO-verify 145) formalise this gap via Breazeal’s taxonomy: PARO occupies the “social interface” level (human-like cues but “shallow models of social cognition”), whereas Sciutti et al.’s (2023, TODO verify: p. 160) vision of robots “knowing what they are doing and why” demands the *socially intelligent* level. The distance between these levels is the cognitive deficit assistive robotics should close.

[] 1.3.2 Medication Adherence and Daily Living Support

Medication non-adherence imposes substantial costs on healthcare systems, and elderly patients with polypharmacy regimens are particularly vulnerable to missed or incorrect doses. Robots in this domain face a different challenge from therapeutic companionship: trust and cognitive load are latent variables that cannot be directly measured, only inferred from noisy behavioural proxies. Lee and See (2004, p. 54) define trust as “the attitude that an agent will help achieve an individual’s goals in a situation characterized by uncertainty and vulnerability”; a definition foregrounding the unobservable nature that necessitates probabilistic modelling. A user may comply with a medication prompt despite low trust (e.g. time pressure), or indeed refuse despite high trust (e.g. task complexity), and thus the observation alone cannot reliably disambiguate the underlying state (Hancock et al., 2011, p. 522).

The Partially Observable Markov Decision Process (POMDP) provides formal machinery for this uncertainty. Chen et al. (2020, p. 6) demonstrate a Trust-POMDP wherein the robot maintains a belief distribution over trust and selects actions that maximise long-term collaboration, showing belief-space planning outperforms fixed strategies in the tested collaborative scenario. Garcez and Lamb (2023, p. 12389) identify the neuro-symbolic paradigm as the ‘third wave’ of AI, wherein neural subsystems (e.g. large language models) handle perception whilst symbolic subsystems (e.g. POMDPs) govern temporal reasoning, providing the temporal scaffold stateless systems lack. Nikolaidis, Hsu and Srinivasa (2017, p. 625) provide empirical corroboration: in a collaborative task ($n = 69$), robots utilising mutual adaptation via a Mixed Observability MDP (modelling human adaptability as a latent variable) were rated significantly more trustworthy than fixed-policy alternatives ($U = 180$, $p = 0.048$). This aligns with Hancock et al.’s (2011, p. 522) finding that robot performance attributes are the strongest trust predictors, whilst demonstrating that belief-space planning as advocated by Chen et al. (2020, p. 6) translates into measurable trust gains.

[] 1.3.3 Physical Rehabilitation and Mobility

Robotic exoskeletons and assistive manipulators for stroke recovery and mobility support constitute a third application domain. These systems need to adapt in real time not only to the patient’s physical state (joint angles and muscle activation patterns) but also to their psychological state: motivation, frustration, and fatigue are internal variables that determine whether a patient perseveres or disengages.

Embodied cognition becomes essential. Brooks (1991, p. TODO) argues that intelligence emerges from physical interaction with the environment rather than abstract representation, whereas Fong, Nourbakhsh and Dautenhahn (2003, p. 149) operationalise this as “perturbatory coupling”: the more channels of mutual influence between robot and environment, the more embodied the system. A rehabilitation robot therefore occupies a uniquely cognitive niche, as it should sense the patient’s body, reason about current capabilities, and adapt appropriately. A purely language-based or screen-based interface cannot achieve this; Matarić et al. (2007, p. TODO) confirm as much empirically, finding that stroke survivors engaged more enthusiastically with a physically embodied assistive robot than with screen-based alternatives. Tapus, Țăpuș and Matarić (2008, p. TODO), in fact show that embodiment alone is insufficient: adaptive personality matching (adjusting interaction distance and speed to the user’s traits) further improved task performance, suggesting rehabilitation robots require not just physical presence but cognitive adaptation to the individual. The cognitive building blocks required (haptic perception, prospective planning of difficulty, episodic memory of the patient’s trajectory) thus suggest an embodied cognitive architecture is necessary rather than a disembodied controller.

□ ### 1.4. Discussion

[] 1.4.1 Challenges

Tapus, Matarić and Scassellati (2007, .pdf-p. 6) projected that by 2012 SAR systems would demonstrate “marked improvement in learning/training/recovery of the user”; yet PARO, the most-deployed platform nearly twenty years later, *still* cannot remember yesterday’s session. Three challenges explain this stalled trajectory. Firstly, computational intractability: solving *POMDPs* exactly is PSPACE-complete (Papadimitriou and Tsitsiklis, 1987, p. 448 {TODO VERIFY}), and the belief simplex grows exponentially with state-space dimensionality. Whilst approximate solvers such as point-based value iteration (Pineau, Gordon and Thrun, 2003, p. 1025; Kaelbling, Littman and Cassandra, 1998, p. 120) and online Monte-Carlo tree search (Silver and Veness, 2010, p. 1) mitigate this, real-time cognitive processing within embodied systems remains an open challenge, particularly when multiple unobserved variables (trust, load, emotion) require tracking simultaneously.

Secondly, the measurement problem: trust, cognitive load, and emotional state are not directly observable; observations thereof are noisy proxies at best. Hancock et al.’s (2011, p. 522) meta-analysis of 29 studies finds that even the strongest correlates of trust explain only modest variance, whilst Broadbent, Stafford and MacDonald (2009, p. TODO) note that acceptance itself depends on matching robot behaviour to user expectations rather than trust alone. Desai et al. (2013, p. 256) further demonstrate that trust dynamics are non-linear, building slowly through consistent performance but degrading rapidly after errors; and thus a single misclassified observation can cascade into inappropriate action selection. Nikolaidis, Hsu and Srinivasa (2017, p. 627), however, demonstrate that mutual adaptation partially mitigates this fragility: when the robot models human adaptability as a latent variable, trust persists even during strategy disagreements, suggesting the variance Hancock et al. report may stem from studies that treat the human as a static rather than co-adaptive partner.

Finally, adaptation without exploitation: a robot that runs inference on cognitive load could, in principle, time its medication requests to coincide with periods of high vulnerability, thereby maximising compliance at the expense of user autonomy. The reward function governing the POMDP’s policy should therefore encode ethical constraints alongside clinicians’ objectives.

[] 1.4.2 Ethical Implications

Assistive robots operating in intimate care spaces (bedrooms, bathrooms, rehabilitation clinics) continuously collect sensitive behavioural data. Facial expressions, vocal patterns, and movement trajectories constitute biometric data, yet regulatory frameworks have not kept pace with deployment. Wachter, Mittelstadt and Floridi (2017, p. TODO) argue that even the General Data Protection Regulation provides no enforceable “right to explanation” of automated decisions; a gap particularly concerning in healthcare wherein recommendations directly affect patient wellbeing.

Moreover, over-reliance on assistive robots risks eroding functional independence. If a robot consistently anticipates and pre-empts needs via prospection, the user may disengage from self-directed activity, thereby creating a dependency that contradicts the assistive mandate. Sharkey (2014, p. 6 (VERIFY)) frames this via Nordenfelt’s ‘Dignity of Identity’: “a robot that dealt impersonally with an older person, without knowing or using their name or their preferences would also be likely to negatively affect their feelings of dignity.” This implies that only cognitively-equipped robots (those with episodic memory of individual users) can avoid dignity violations; reactive systems such as PARO, regardless of their therapeutic benefits (Wada and Shibata, 2007, p. 974), risk infantilisation precisely because they cannot personalise. The responsibility gap compounds this further: when a care robot administers incorrect medication, liability falls ambiguously between manufacturer, deployer, and clinician.

Furthermore, the deployment of assistive robots is not equitable: wealthy nations with sufficient infrastructure and research investment stand to benefit, whilst low-income populations face a widening digital divide in access to care technologies. The question of whether assistive robots displace human carers or augment them remains unresolved.

The ethical watchword is therefore proactive regulation: design-stage ethics that anticipate failure modes before deployment, rather than reactive patchwork after harm. Per the embodied cognition thesis, if intel-

ligence indeed requires a body, and that body enters the most intimate spaces of vulnerable persons, then the ethical stakes of assistive cognitive robotics are uniquely high.

□ ## 1.5. Conclusion

Assistive robotics stands at an inflection point. Current systems (PARO, medication prompt robots, rehabilitation aids) deliver measurable benefits within narrow operational envelopes, yet their reactive architectures limit sustained, personalised effectiveness. The Vernon, Metta and Sandini (2007) cognition cycle provides the architectural blueprint for graduating beyond this plateau: assistive robots that anticipate (prospection), remember (episodic and semantic memory), reason about others’ mental states (theory of mind), and monitor their own performance (metacognition) would constitute a qualitative advance over the most-capable systems deployed.

The neuro-symbolic paradigm offers a viable path toward this vision, as the Trust-POMDP framework attests (Chen et al., 2020). Sciutti et al. (2023, pp. 162-163) independently identify the integration of learning with model-based approaches as cognitive robotics’ most-prominent trajectory; that this converges with Garcez and Lamb’s (2023, p. 12389) ‘third wave’ thesis from AI theory suggests the direction is robust rather than parochial. Future applications will likely extend beyond single-task assistance toward cognitively autonomous home-dwelling companions: robots that proactively monitor health indicators, anticipate daily needs via episodic memory, and adapt their interaction style to the user’s evolving cognitive and emotional state. Sharkey and Sharkey (2012, p. 27) identify this trajectory whilst cautioning that such systems risk replacing rather than supplementing human-care, and therefore the field should pursue cognitive capability and ethical governance in concert, lest the technology displace these very carers it was meant to supplement. Figure-2 visualises this trajectory. The ultimate test, per the embodied cognition thesis, is a robot that can: sense → remember → anticipate → adapt within the physical world, whilst respecting the autonomy, and dignity, of the persons it serves.




Figure 2: Trajectory of assistive robotics from reactive single-task systems (PARO) through adaptive belief-based architectures (Trust-POMDP) toward cognitively autonomous home-dwelling companions. Expanding assistive scope without expanding cognitive capability is insufficient; the diagonal trajectory requires both.

1.6. Task-3 References

1.6.1 Papers and Journal Articles

- ☐ verify all links work and extract each papers insights

REMOVE: one checkmark for working link second for insight gathered

- ☐ [] Broadbent, E., Stafford, R. and MacDonald, B. (2009) ‘Acceptance of Healthcare Robots for the Older Population: Review and Future Directions’, *International Journal of Social Robotics*, 1(4), pp. 319-330. Available at: https://www.researchgate.net/publication/220397395_Acceptance_of_Healthcare_Robots (Accessed: 25 March 2026).
- ☐ [] Brooks, R. A. (1991) ‘Intelligence without representation’, *Artificial Intelligence*, 47(1-3), pp. 139-159. Available at: <https://people.csail.mit.edu/brooks/papers/representation.pdf> (Accessed: 24 March 2026).
- ☐ [] Chen, M., Nikolaidis, S., Soh, H., Hsu, D. and Srinivasa, S. (2020) ‘Trust-Aware Decision Making for Human-Robot Collaboration: Model Learning and Planning’, *ACM Transactions on Human-Robot Interaction*, 9(2), pp. 1-23. Available at: <https://arxiv.org/abs/1801.04099> (Accessed: 15 March 2026).
- ☐ [] Desai, M., Kaniarasu, P., Medvedev, M., Steinfeld, A. and Yanco, H. (2013) ‘Impact of robot failures and feedback on real-time trust’, *Journal of Human-Robot Interaction*, 2(1), pp. 251-275. Available at: <https://ieeexplore.ieee.org/document/6483596> (Accessed: 20 March 2026).
- ☐ [] Fong, T., Nourbakhsh, I. and Dautenhahn, K. (2003) ‘A survey of socially interactive robots’, *Robotics and Autonomous Systems*, 42(3-4), pp. 143-166. Available at: <https://www.cs.cmu.edu/~illah/PAPERS/socialr> (Accessed: 18 March 2026).
- ☐ [] Garcez, A. d’A. and Lamb, L. C. (2023) ‘Neurosymbolic AI: The 3rd Wave’, *Artificial Intelligence Review*, 56, pp. 12387-12406. Available at: <https://link.springer.com/article/10.1007/s10462-023-10448-w> (Accessed: 20 March 2026).
- ☐ [] Hancock, P. A., Billings, D. R., Schaefer, K. E., Chen, J. Y. C., de Visser, E. J. and Parasuraman, R. (2011) ‘A meta-analysis of factors affecting trust in human-robot interaction’, *Human Factors*, 53(5), pp. 517-527. Available at: <https://journals.sagepub.com/doi/10.1177/0018720811417254> (Accessed: 15 March 2026).
- ☐ [] Kaelbling, L. P., Littman, M. L. and Cassandra, A. R. (1998) ‘Planning and acting in partially observable stochastic domains’, *Artificial Intelligence*, 101(1-2), pp. 99-134. Available at: <https://people.csail.mit.edu/lpk/papers/aij98-pomdp.pdf> (Accessed: 13 March 2026).
- ☐ [] Laird, J. E. (2012) *The Soar Cognitive Architecture*. Cambridge, MA: MIT Press.
- ☐ [] Lee, J. D. and See, K. A. (2004) ‘Trust in automation: Designing for appropriate reliance’, *Human Factors*, 46(1), pp. 50-80. Available at: https://journals.sagepub.com/doi/10.1518/hfes.46.1.50_30392 (Accessed: 15 March 2026).
- ☐ [] Matarić, M. J., Eriksson, J., Feil-Seifer, D. J. and Winstein, C. J. (2007) ‘Socially assistive robotics for post-stroke rehabilitation’, *Journal of NeuroEngineering and Rehabilitation*, 4(5), pp. 1-9. Available at: <https://pmc.ncbi.nlm.nih.gov/articles/PMC1821334/> (Accessed: 25 March 2026).
- ☐ [] Nikolaidis, S., Hsu, D. and Srinivasa, S. (2017) ‘Human-robot mutual adaptation in collaborative tasks: Models and experiments’, *The International Journal of Robotics Research*, 36(5-7), pp. 618-634. Available at: <https://journals.sagepub.com/doi/10.1177/0278364917690593> (Accessed: 20 March 2026).
- ☐ [] Papadimitriou, C. H. and Tsitsiklis, J. N. (1987) ‘The complexity of Markov decision processes’, *Mathematics of Operations Research*, 12(3), pp. 441-450. Available at: <https://web.mit.edu/jnt/www/Papers/J016-87-mdp-complexity.pdf> (Accessed: 13 March 2026).
- ☐ [] Pineau, J., Gordon, G. and Thrun, S. (2003) ‘Point-based value iteration: An anytime algorithm for POMDPs’, in *Proceedings of the 18th International Joint Conference on Artificial Intelligence (IJCAI-03)*, pp. 1025-1030. Available at: <http://www.cs.cmu.edu/~ggordon/jpineau-ggordon-thrun.ijcai03.pdf> (Accessed: 24 March 2026).
- ☒ [] Sciutti, A., Beetz, M., Inamura, T., Korsah, A., Oh, J., Sandini, G., Shimoda, S. and Vernon, D. (2023) ‘The Present and the Future of Cognitive Robotics’, *IEEE Robotics & Automation Magazine*, 30(3), pp. 160-163. Available at: <https://ieeexplore-ieee-org.plymouth.idm.oclc.org/document/10255092> (Accessed: 18 March 2026).

- [] Sharkey, A. (2014) ‘Robots and human dignity: A consideration of the effects of robot care on the dignity of older people’, *Ethics and Information Technology*, 16(1), pp. 63-75. Available at: <https://philarchive.org/rec/SHARAH-2> (Accessed: 22 March 2026).
- [] Sharkey, A. and Sharkey, N. (2012) ‘Granny and the robots: ethical issues in robot care for the elderly’, *Ethics and Information Technology*, 14(1), pp. 27-40. Available at: <https://philarchive.org/rec/SHAGAT> (Accessed: 22 March 2026).
- [] Silver, D. and Veness, J. (2010) ‘Monte-Carlo planning in large POMDPs’, in *Advances in Neural Information Processing Systems (NeurIPS 23)*, pp. 2164-2172. Available at: <https://proceedings.neurips.cc/paper/2010/file/Paper.pdf> (Accessed: 24 March 2026).
- [] Tapus, A., Matarić, M. J. and Scassellati, B. (2007) ‘Socially assistive robotics [Grand Challenges of Robotics]’, *IEEE Robotics & Automation Magazine*, 14(1), pp. 35-42. Available at: <https://scazlab.yale.edu/sites/default/files/files/Tapus-RAM2007.pdf> (Accessed: 25 March 2026).
- [] Tapus, A., Țăpuș, C. and Matarić, M. J. (2008) ‘User-robot personality matching and assistive robot behavior adaptation for post-stroke rehabilitation therapy’, *Intelligent Service Robotics*, 1(2), pp. 169-183. Available at: <https://hal.science/hal-00770108/document> (Accessed: 26 March 2026).
- ☒ [] Vernon, D., Metta, G. and Sandini, G. (2007) ‘A Survey of Artificial Cognitive Systems: Implications for the Autonomous Development of Mental Capabilities in Computational Agents’, *IEEE Transactions on Evolutionary Computation*, 11(2), pp. 151-180. Available at: https://www.robotcub.org/misc/papers/07_Vernon_Metta_Sandini_IEEE.pdf (Accessed: 13 March 2026).
- [] Wachter, S., Mittelstadt, B. and Floridi, L. (2017) ‘Why a Right to Explanation of Automated Decision-Making Does Not Exist in the General Data Protection Regulation’, *International Data Privacy Law*, 7(2), pp. 76-99. Available at: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=2903469 (Accessed: 22 March 2026).
- [] Wada, K. and Shibata, T. (2007) ‘Living with seal robots: its sociopsychological and physiological influences on the elderly at a care house’, *IEEE Transactions on Robotics*, 23(5), pp. 972-980. Available at: <https://ieeexplore.ieee.org/document/4339551> (Accessed: 18 March 2026).

2- Task (4) Novel Programming Project

- CRITICAL: reconfigure report below to match newest gaze.py
- ☒ multi-layer defence-in-depth: talk about how silero-vad was used to stop whisper hallucinations
- verify word count

2.1. Introduction (10%; Salman)




Figure 3: GAZE system architecture. Five live input signals feed the AdaptiveEngine (the symbolic reasoning layer), wherein a context block is built and passed to GPT-5.4 function-calling; a separate deterministic GPT-4.1 call handles answer verification (i.e. the referee, not the host). The generated response delivers concurrently via Pepper’s speech, gesture, and LED subsystems, whilst the dashboard pulls frames from Pepper’s ALVideoDevice socket stream rather than still photos. The loop circulates to the next round, now adapted.

As noted in {video of easy questions} wherein Salman purposely disengages from the human-robot interaction, and GAZE-Pepper

2.2. Background (10%; Alfie)

GAZE sits within socially assistive robotics (*the deployment of robots to support users through social interaction rather than physical contact*) and affective computing, which Picard (1997, p. 1, Abstract) founds as “computing that relates to, arises from, or influences emotions”; Tapus, Matarić and Scassellati (2007, p. 35) define this sub-field as systems that “assist users through social interaction.” Most-current platforms react to a single input signal (Fong, Nourbakhsh and Dautenhahn, 2003, p. 148), suffering the single-signal problem: a facial-expression classifier misreads resting faces as displeasure; a response-time metric mistakes deliberation for disengagement. Calvo and D’Mello (2010, p. 28) identify “the inherent challenges with unisensory affect detection”; Poria et al. (2017, p. 99) report that multimodal systems were “consistently (85% of systems) more accurate than their best unimodal counterparts, with an average improvement of 9.83%.” Spezialetti, Placidi and Rossi (2020, pp. 1-2, Introduction) corroborate this within HRI specifically, reviewing recognition systems across facial, vocal, and physiological channels. This demands multi-signal fusion (the system weighs complementary channels together instead of trusting any one alone).

- PROOFREAD GAZE’s core contribution is multi-signal emotional inference: facial expression (CNN, Workshop 10), vocal emotion (MLP, Workshop 8), speech volume/RMS, response time, answer correctness, and answer text fuse with derived temporal signals into a single-inferred userstate. The implementation is face-primary: voice nudges only when the face is Neutral and the MLP clears 0.9 confidence (**fearful** excluded as the silence-attractor). GPT-5.4 generates dialogue whilst the symbolic AdaptiveEngine governs state inference, grounding output in Pepper’s affordances (Ahn et al., 2022, p. 1); a GPT-4.1 verifier handles answer-checking. Smedegaard (2019, p. 4) warns engagement reflects novelty rather than sustained interest; GAZE targets this via adaptive... TODO finish sentence (game-switching?)

2.3. Methods & Setup (35%; Alfie)

2.3.1 System Architecture

- PROOFREAD GAZE operates as a conversational loop rather than a rigid question-answer cycle. Each turn: 1) Pepper’s streamed camera frame (live `ALVideoDevice` socket, not per-turn photo) is classified for facial expression; 2) Pepper records through all four microphones with ambient-calibrated silence detection; 3) the WAV is canonicalised to mono 16 kHz by loudest-channel selection, then processed for vocal emotion, volume/RMS, and Whisper transcription; 4) the AdaptiveEngine infers user-state from face, voice, volume, response time, correctness, and answer text; 5) context plus user speech are sent to GPT-5.4 with four function-calling tools; 6) speech, gesture, and LED state fire concurrently on Pepper. Computation runs on the laptop; Pepper handles physical I/O. The dashboard hence reflects live emotional-state inference from session start.
- This function-calling architecture is neuro-symbolic: GPT-5.4 governs dialogue and decision-making, whilst the AdaptiveEngine and game logic are exposed as callable tools. This aligns with Garcez and Lamb’s (2023, p. TODO: VERIFY 12389) ‘third wave’ paradigm, wherein neural and symbolic components share a structured interface (cf. Ahn et al., 2022, p. 1).

essentially in many cases the robot will output text which controll descisiom-making (if says ‘x’ do ‘y’)

2.3.2 Input Layer: Multimodal and Behavioural Signals

1- Facial Expression (vision-based). A pre-trained CNN (Workshop 10) classifies the user’s expression into seven categories (Angry, Disgust, Fear, Happy, Neutral, Sad, Surprise) from a 48×48 greyscale face region, building upon Ekman and Friesen (1971, pp. 127-128), whose cross-cultural results show that “particular facial behaviors are universally associated with particular emotions,” finding that even preliterate (without written language) cultures with “minimal opportunity to have learned to recognize uniquely Western facial expressions” identified the same six emotions; this taxonomy remains “still the most popular

perspective for FER” (Li and Deng, 2020, p. 1). Known limitations (cultural bias, resting-face misclassification) are mitigated by combining facial evidence with behavioural signals (response time, correctness, silence history, volume/RMS) rather than allowing the CNN to decide user state alone.

- **PROOFREAD 2- Verbal Answer (speech-based).** Pepper records through all four microphones (the [1,1,1,1] mask) via `ALAudioRecorder` at 16 kHz, whilst `ALAudioDevice` polls the maximum energy across all four mics, hence catching off-axis speakers a front-only poll would miss. Recording terminates after 1.2 seconds of post-speech silence or at the adaptive hard ceiling. After SFTP, `force_mono_16k_wav()` selects the loudest channel by RMS and rewrites mono 16 kHz PCM. The canonical WAV is then transcribed via OpenAI Whisper (`whisper-1`), whose models “approach [human] accuracy and robustness” (Radford et al., 2023, Abstract). Whisper is utilised as a delivery-aware sensor: `verbose_json` exposes per-segment `no_speech_prob`, `avg_logprob`, and `compression_ratio` diagnostics; it therefore encodes *how* the user spoke alongside *what* they said cueing older ASR strips wholesale. Turn-taking is strictly sequential, wherein the mic disengages during TTS playback; barge-in (user-initiated interrupt mid-utterance) was therefore deliberately omitted, owing to NAOqi’s imperfect acoustic echo cancellation (which would have the robot’s own voice fire false interrupts) and the resultant false-trigger risk under noisy demo conditions.

Listing 1: `nao_record`: calibrated-threshold silence detection polling the max energy across all four Pepper mics, with firmware fallback. Recording uses the [1,1,1,1] channel mask; `force_mono_16k_wav()` downstream picks the loudest channel by RMS so Whisper, Silero-VAD, the WS-08 MLP, and `measure_volume()` all see one canonical mono 16 kHz layout.

```

1 def nao_record(ssh, energy_threshold: int = DEFAULT_ENERGY_THRESHOLD,
2               record_max_secs: float = RECORD_MAX_SECS,
3               silence_secs: float = SILENCE_DURATION):
4     """Record audio on Pepper with dynamic silence detection.
5     - poll the max energy across all four mics to stop recording early if silence detected
6     - calibrated energy threshold to avoid false positives from ambient noise
7     - if getFrontMicEnergy is unsupported (e.g. older firmware), fall back to a safe
      ↪ fixed-duration recording to ensure the demo still works, albeit without silence
      ↪ detection
8     """
9
10    nao_run(ssh, f"""
11 from naoqi import ALProxy
12 import time
13
14 rec = ALProxy("ALAudioRecorder", "127.0.0.1", 9559)
15
16 rec.stopMicrophonesRecording()
17 rec.startMicrophonesRecording("{REMOTE_WAV}", "wav", 16000, [1, 1, 1, 1])
18
19 try:
20     audio = ALProxy("ALAudioDevice", "127.0.0.1", 9559)
21
22     speech_detected = False
23     silence_start = None
24     start = time.time()
25     threshold = {energy_threshold}
26
27     while True:
28         elapsed = time.time() - start
29
30         # hard ceiling; never exceed max duration
31         if elapsed >= {record_max_secs}:
32             break
33
34         # poll all four mics and take the MAX; a user stood to either side of Pepper registers on the

```

```

    ↪ side mics but not the front, so front-only polling misses them and the silence-detector
    ↪ stops recording mid-utterance
35     try:
36         energy = max(
37             audio.getFrontMicEnergy(),
38             audio.getLeftMicEnergy(),
39             audio.getRightMicEnergy(),
40             audio.getRearMicEnergy(),
41         )
42     except Exception:
43         energy = audio.getFrontMicEnergy() # firmware fallback
44
45     if elapsed < {RECORD_MIN_SECS}:
46         # minimum recording period
47         if energy > threshold:
48             speech_detected = True
49             time.sleep({SILENCE_POLL_SECS})
50             continue
51
52     if energy > threshold:
53         speech_detected = True
54         silence_start = None
55     else:
56         if speech_detected and silence_start is None:
57             silence_start = time.time()
58         if speech_detected and silence_start is not None:
59             if (time.time() - silence_start) >= {silence_secs}:
60                 break
61
62     time.sleep({SILENCE_POLL_SECS})
63
64 except Exception as e:
65     # firmware fallback; getFrontMicEnergy() unsupported on this Pepper
66     # fall back to a safe fixed-duration recording so the demo never breaks
67     print(" [Silence detection failed: " + str(e) + "] Falling back to fixed-duration recording")
68     time.sleep({record_max_secs})
69
70 rec.stopMicrophonesRecording()
71 """
72 sftp = ssh.open_sftp()
73 sftp.get(REMOTE_WAV, LOCAL_WAV)
74 sftp.close()

```

- PROOFREAD Whisper hallucinates on near-silent audio, emitting training-set tics (CJK gibberish, prompt-primed repetition); the latter NLG *degeneration* wherein models get “stuck in repetitive loops” (Ji et al., 2023, p. 3). Transcription is gated by five layers: 1) a near-empty-file RMS floor (NAO_MIN_RMS_TO_TRANSCRIBE = 30, emergency only; the older 120 wrongly rejected short answers like “yes” or “Tom”); 2) Silero-VAD, applied to both LOCAL and NAO modes since `force_mono_16k_wav()` canonicalises beforehand; 3) Vosk grammar-restricted wake-word gate (“Pepper”/“Gaze”), bypassable for the first-turn name prompt; 4) Whisper’s own `no_speech_prob` / `avg_logprob` / `compression_ratio` signals; 5) an aggressively-normalised hallucination blacklist.

Listing 2: `transcribe`: five-layer defence chain against Whisper hallucination on near-silent audio (`gaze.py`, lines 1357–1427).

```

1 def transcribe(bypass_wake_word: bool = False,
2               whisper_prompt: str = "User answers a quiz or chats with a companion robot.") -> str:
3     """
4     Transcribe the local WAV with Whisper via a five-layer defence:

```

```

5     1- existing RMS + speech_detected pre-gate (in conversation_loop)
6     2- Silero VAD hard gate (this function; both LOCAL and NAO modes)
7     3- Vosk wake-word gate (requires "Pepper"/"Gaze"; bypassable)
8     4- Whisper's own no_speech_prob + avg_logprob (from verbose_json)
9     5- aggressively-normalised hallucination blacklist
10
11 Returns "" on failure, on any gate rejecting, or when the
12 transcription normalises to a known-hallucination phrase.
13
14 bypass_wake_word=True disables Layer 3 for turns where no wake-word
15 is expected (e.g. the first-turn name prompt).
16 ""
17 if LOCAL_MODE and not _local_speech_detected:
18     return ""
19
20 # Layer 2: Silero VAD hard gate; applied to both LOCAL and NAO modes because
21 ↳ force_mono_16k_wav() canonicalises the WAV before transcribe() runs
22 if not has_real_speech(LOCAL_WAV):
23     print(" Silero VAD found no speech; skipping Whisper.")
24     return ""
25
26 # Layer 3: Vosk wake-word gate. User must say "Pepper"/"Gaze"; bypassable for the first-turn
27 ↳ name prompt where no wake-word is expected.
28 if not bypass_wake_word and not has_wake_word(LOCAL_WAV):
29     print(" No wake-word detected; skipping Whisper.")
30     return ""
31
32 try:
33     with open(LOCAL_WAV, "rb") as fh:
34         resp = client.audio.transcriptions.create(
35             model="whisper-1",
36             file=fh,
37             response_format="verbose_json",
38             temperature=0.0,
39             prompt=whisper_prompt,
40             timeout=API_TIMEOUT,
41         )
42     text = (getattr(resp, "text", "") or "").strip()
43
44 # Layer 4: Whisper self-signals (no_speech_prob + avg_logprob + compression_ratio per
45 ↳ segment). resp is a Pydantic TranscriptionVerbose thus attribute access (not dict
46 ↳ indexing).
47 segments = getattr(resp, "segments", None) or []
48 if segments:
49     no_speech_vals = [s.no_speech_prob for s in segments
50                       if getattr(s, "no_speech_prob", None) is not None]
51     logprob_vals = [s.avg_logprob for s in segments
52                    if getattr(s, "avg_logprob", None) is not None]
53     compression_vals = [s.compression_ratio for s in segments
54                         if getattr(s, "compression_ratio", None) is not None]
55 if no_speech_vals and max(no_speech_vals) > 0.6:
56     print(f" Whisper flagged silence (max no_speech_prob="
57           f"{max(no_speech_vals):.2f}); dropping {text!r}")
58     return ""
59 if logprob_vals and min(logprob_vals) < -1.0:
60     print(f" Whisper low-confidence (min avg_logprob="
61           f"{min(logprob_vals):.2f}); dropping {text!r}")
62     return ""
63 # compression_ratio > 2.4 catches repetition loops (e.g. "Q1. Q1. Q1..."); repeating

```

```

60         ↪ tokens compress well. 2.4 matches local Whisper's default.
61     if compression_vals and max(compression_vals) > 2.4:
62         print(f" Whisper repetition loop (max compression_ratio="
63             f"{max(compression_vals):.2f}); dropping {text!r}")
64         return ""
65     # Layer 5: aggressively-normalised hallucination blacklist
66     if is_known_hallucination(text):
67         print(f" Filtered Whisper hallucination: {text!r}")
68         return ""
69
70     # Strip leading "Pepper"/"Gaze" so handlers receive just the answer; \b blocks
71     ↪ "Pepperoni"/"Gazebo" false positives..
72     stripped = re.sub(r'(?i)^\s*(pepper|gaze)\b[.\s]*', '', text).strip()
73     return stripped
74 except Exception as e:
75     print(f" Whisper transcribe failed ({e}); returning empty")
76     return ""

```

3- Vocal Emotion (audio-based). The same WAV is passed through a pre-trained MLP (Workshop 8) *before* transcription, classifying vocal state into four emotions (calm, happy, fearful, disgust) via MFCC, chroma, and mel-spectrogram features; El Ayadi, Kamel and Karray (2011, p. 577) identify MFCCs as “the most promising features” for speech-emotion recognition. This provides a second, independent modality; the two may disagree, wherein decision logic arbitrates.

Listing 3: `SpeechEmotionModel.extract_features`: MFCC + chroma + mel-spectrogram feature vector matching the WS-08 training pipeline, with peak-normalisation for quieter brain-injured users (`gaze.py`, lines 192–220).

```

1
2 # PASTE SpeechEmotionModel.extract_features FROM gaze.py HERE (lines ~192--220, including the
   ↪ @staticmethod decorator)

```

4- Response Time (engagement-based). A Python timer measures elapsed time from question delivery to recording completion; the Whisper call occurs *after* the timer halts, isolating deliberation time from API latency.

5- Answer Correctness (task-based). GPT-4.1’s `check_game_answer` evaluates the transcribed answer at a deterministic temperature; the resultant binary correctness feeds the AdaptiveEngine and the derived rolling-correctness signal.

2.3.3 Process Layer: Multi-Signal State Inference

- PROOFREAD The AdaptiveEngine’s `infer_state()` combines multimodal and behavioural signals to classify the user into one of five states: *Thriving*, *Comfortable*, *Struggling*, *Frustrated*, or *Disengaged*. Five main input channels are supplemented by three derived temporal features: rolling correctness over the last five rounds, consecutive-silence count, and consecutive-wrong streak length. Crucially, classification is face-primary: facial expression, correctness, response time, silence, wrong-streaks, and volume carry the decision; vocal emotion fires only as a high-confidence tie-breaker when face is Neutral AND the MLP clears 0.9 AND the label is not **fearful** (the silence-attractor). Voice is thus retained as evidence but never allowed to override stronger behavioural readings, addressing the brittleness Desai et al. (2013, p. 256) observe in single-signal systems. Thresholds (correctness floor 0.4, ceiling 0.8, response-time baseline 15s, consecutive-wrong trigger 3, silence threshold 2) were derived from pilot testing.




Figure 4: Multi-signal inference pipeline. Five raw inputs captured each round and three derived temporal signals computed from session history feed into `infer_state()`, which applies cross-modal rules (e.g. camera reads *Angry* but user answers fast and correctly $\rightarrow$ *Comfortable*) to classify the user into one of five states. Voice is treated as advisory rather than dominant: it nudges the state only when the face is Neutral AND the MLP’s confidence clears 0.9 AND the label is not *fearful* (the silence-attractor), hence a noisy vocal label can never override stronger visual or behavioural evidence; therein lies the multi-signal novelty.

The complete classification logic, face-primary with voice tie-breakers, is shown in Listing~4 (extracted from `gaze.py infer_state()`):

Listing 4: Multi-signal inference rules (verbatim excerpt from `gaze.py infer_state()`; method-body indentation stripped). The newest implementation is face-primary: every rule fires off facial expression + behavioural signals first, voice is consulted only as a high-confidence tie-breaker at the end.

```

1
2 # 1- FACE-PRIMARY RULES: these fire before voice is ever consulted
3
4 # thriving: performing well + fast responses
5
6 if (correctness >= CORRECTNESS_CEILING and response_time < RESPONSE_TIME_BASELINE * 0.5):
7     return InferredState.THRIVING
8 if expression == "Angry" and correct and response_time < RESPONSE_TIME_BASELINE * 0.6:
9     return InferredState.COMFORTABLE
10
11 # disengaged: silence threshold, or face + slow response + poor performance
12
13 if self.consecutive_silences >= SILENCE_THRESHOLD:
14     return InferredState.DISENGAGED
15 if (expression == "Neutral"
16     and response_time > RESPONSE_TIME_BASELINE
17     and correctness < 0.5):
18     return InferredState.DISENGAGED
19 if (low_arousal and expression == "Neutral"
20     and response_time > RESPONSE_TIME_BASELINE * 0.8):
21     return InferredState.DISENGAGED
22
23 # frustrated: negative face + poor performance, or 3+ wrong + negative face, or high arousal +
    ↪ negative face + poor

```

```

24
25 if expression in ("Angry", "Disgust") and correctness < CORRECTNESS_FLOOR:
26     return InferredState.FRUSTRATED
27 if self.consecutive_wrong >= 3 and expression in ("Angry", "Sad", "Fear"):
28     return InferredState.FRUSTRATED
29 if (high_arousal and expression in ("Angry", "Disgust", "Fear")
30     and correctness < CORRECTNESS_FLOOR):
31     return InferredState.FRUSTRATED
32
33 # struggling: sadness + slow; poor correctness; fear + wrong
34
35 if expression == "Sad" and response_time > RESPONSE_TIME_BASELINE * 0.7:
36     return InferredState.STRUGGLING
37 if correctness < CORRECTNESS_FLOOR:
38     return InferredState.STRUGGLING
39 if expression == "Fear" and not correct:
40     return InferredState.STRUGGLING
41
42 # 2- VOICE TIE-BREAKERS: only when face is Neutral AND voice is trustworthy. Voice can nudge but
43     ↪ never override a face-driven verdict.
44
45 if expression == "Neutral" and trust_voice:
46     if vocal_emotion == "happy" and correct and correctness >= CORRECTNESS_CEILING:
47         return InferredState.THRIVING
48     if vocal_emotion == "calm" and correctness >= 0.5:
49         return InferredState.COMFORTABLE
50
51 return InferredState.COMFORTABLE

```

The `decide()` function then maps the inferred state to concrete adaptive actions: difficulty adjustment (easy, medium, hard), game switching (numbers ↔ letters when the user is frustrated or disengaged), hint provision, encouragement, and tone selection. This parallels Nikolaidis, Hsu and Srinivasa’s (2017, p. 625) mutual-adaptation paradigm. Table~1 summarises the mapping.

Inferred State	Difficulty	Game Switch	Hint	Tone	LED Colour
Thriving	↑ ramp up	No	No	Energetic	Green
Comfortable	↑ if >70%	No	No	Neutral	Cyan
Struggling	↓ ease off	No	Yes	Encouraging	Yellow
Frustrated	↓↓ Easy	After 4 wrong	No	Calm	Red
Disengaged	—	After 3 checkouts	No	Energetic	Magenta

Table 1: State-to-action mapping. The adaptive engine translates each inferred state into a specific combination of difficulty adjustment, game-type switch, hint provision, tone selection, and LED colour. Arrows indicate direction of difficulty change relative to the current level, wherein the mapping applies solely to `decide()`’s engine output (the spoken-checkout path, e.g. repeated *skip*); the main-loop tiered intervention on pure-silence turns, tier-1 check-in at 3 silences and tier-2 switch at 4, bypasses `decide()` entirely.

Disengagement is flagged by three OR’d rules wherein any one trips it: 1) two consecutive silences (mic-empty or “skip”/“pass” checkout phrases), 2) neutral-face + response > 15s + rolling-correctness < 50% (the cognitive-checkout pattern), 3) quiet-voice + neutral-face + slow-ish response (early-warning drift, catches it before accuracy tanks). GAZE then tiers its intervention, gentle check-in at three silences, game-switch at four, wherein each tier fires at most once per silent spell.

2.3.4 Generate Layer: Dynamic Prompt Construction

- PROOFREAD
- Rather than constructing a fixed game prompt each round, GAZE utilises OpenAI function calling to let GPT-5.4 decide actions. Every turn, `build_signal_context()` packages live signals (face, voice, volume, response time, rolling accuracy, recent-face/recent-vocal windows) into a context block prepended to the user’s transcribed speech. This is sent to GPT-5.4 with four callable tools: `generate_game_question`, `check_game_answer`, `evaluate_last_adaptation`, and `request_more_time`. The LLM decides which to invoke; during natural conversation it calls none, whilst during gameplay it sequences `check_game_answer` and `generate_game_question`. Game-question generation runs as a separate JSON-only call with a variety seed plus a rolling 30-question do-not-repeat memory, hence preventing mode-collapsed targets. Answer verification is delegated to a deterministic GPT-4.1 verifier (temperature 0.0), reducing the hallucination risk Ji et al. (2023, p. 3) identify. A 10-second timeout wraps every API call in case the network stalls so Pepper falls back gracefully rather than freezing mid-turn.

2.3.5 Output Layer: Aligned Multimodal Response

Speech, gestures, and LED state fire concurrently via threading. Context-aligned gestures (e.g. animated speech for celebratory moments) execute alongside dialogue, whilst LED colours reflect the inferred state, providing a secondary non-verbal feedback channel.

2.3.6 Adaptation Self-Evaluation and Session Persistence

- PROOFREAD GAZE exposes `evaluate_adaptation()` to the LLM as the `evaluate_last_adaptation` tool, hence the LLM can compare consecutive rounds to judge whether a previous adaptation helped (e.g. did a difficulty decrease after frustration produce a correct answer?). Session progress is saved to `game_save.json` after every game round, with a belt-and-braces auto-save every two turns regardless, in case the session is interrupted before the next scheduled save.

2.4. Outcome & System Analysis (30%; Salman)

2.5. Conclusion (10%; Alfie)

- PROOFREAD GAZE implements multi-signal emotional inference across two independent modalities (*facial expression via WS-10 CNN and vocal emotion via WS-08 MLP*) alongside speech volume/RMS, response time, answer correctness, answer text, and derived temporal signals, hence yielding a more robust user-state estimate than any single channel. The implementation hardens further by recording all four Pepper microphones, canonicalising audio to mono 16 kHz by loudest-channel selection, applying Silero-VAD to both modes, and streaming Pepper’s camera via persistent `ALVideoDevice` rather than per-turn SFTP. The hybrid architecture (GPT-5.4 dialogue paired with symbolic rule-based inference and a deterministic GPT-4.1 verifier) and adaptive game-switching directly target the novelty-decay problem Smedegaard (2019, p. 4) identifies.
- PROOFREAD The conversational architecture, wherein the LLM decides actions via function calling, positions GAZE as a social companion rather than a rigid game host. <!-- WRAPPED OUT: factually
→ inaccurate (no personality-mode selection in current gaze.py)The selectable personality
→ modes extend Tapus, Tapus and Mataric's (2008, p. TODO: VERIFY PAGE)personality matching
→ from static trait-matching to dynamic, signal-driven adaptation. -->

The multi-signal approach could transfer to stroke rehabilitation re-engagement (Mataric et al., 2007), educational tutoring, or neurodivergent support. Future work could replace hand-coded thresholds with learned parameters from longitudinal data, and fine-tune both models on in-session Pepper captures.

2.6 Task-4 References (5%)

- verify all references

Alfie's

- Ahn, M., Brohan, A., Brown, N., et al. (2022) ‘Do As I Can, Not As I Say: Grounding Language in Robotic Affordances’, *arXiv preprint arXiv:2204.01691*. Available at: <https://arxiv.org/abs/2204.01691> (Accessed: 24 March 2026).
- [X]
- [] Desai, M., Kaniarasu, P., Medvedev, M., Steinfeld, A. and Yanco, H. (2013) ‘Impact of robot failures and feedback on real-time trust’, *Journal of Human-Robot Interaction*, 2(1), pp. 251–275. Available at: <https://ieeexplore.ieee.org/document/6483596> (Accessed: 20 March 2026).
- [X]
- [X]
- [] Fong, T., Nourbakhsh, I. and Dautenhahn, K. (2003) ‘A survey of socially interactive robots’, *Robotics and Autonomous Systems*, 42(3–4), pp. 143–166. Available at: <https://www.cs.cmu.edu/~illah/PAPERS/socialroboticssurvey.pdf> (Accessed: 18 March 2026).
- [] Garcez, A. d’A. and Lamb, L. C. (2023) ‘Neurosymbolic AI: The 3rd Wave’, *Artificial Intelligence Review*, 56, pp. 12387–12406. Available at: <https://link.springer.com/article/10.1007/s10462-023-10448-w> (Accessed: 20 March 2026).
- [] Ji, Z., Lee, N., Frieske, R., et al. (2023) ‘Survey of Hallucination in Natural Language Generation’, *ACM Computing Surveys*, 55(12), pp. 1–38. Available at: <https://dl.acm.org/doi/10.1145/3571730> (Accessed: 22 March 2026).
- [] Kaelbling, L. P., Littman, M. L. and Cassandra, A. R. (1998) ‘Planning and acting in partially observable stochastic domains’, *Artificial Intelligence*, 101(1–2), pp. 99–134. Available at: <https://people.csail.mit.edu/lpk/papers/aij98-pomdp.pdf> (Accessed: 13 March 2026).
- [X]
- [] Nikolaidis, S., Hsu, D. and Srinivasa, S. (2017) ‘Human-robot mutual adaptation in collaborative tasks: Models and experiments’, *The International Journal of Robotics Research*, 36(5–7), pp. 618–634. Available at: <https://journals.sagepub.com/doi/10.1177/0278364917690593> (Accessed: 20 March 2026).
- [] Picard, R.W. (1997) ‘Affective Computing’, *MIT Media Laboratory Perceptual Computing Section Technical Report No. 321*. Available at: <https://affect.media.mit.edu/pdfs/95.picard.pdf> (Accessed: 14 April 2026).
- [X]
- [] Radford, A., Kim, J.W., Xu, T., Brockman, G., McLeavey, C. and Sutskever, I. (2023) ‘Robust Speech Recognition via Large-Scale Weak Supervision’, in *Proceedings of the 40th International Conference on Machine Learning (ICML 2023)*, PMLR 202, pp. 28492–28518. Available at: <https://proceedings.mlr.press/v202/radford23a/radford23a.pdf> Proceedings of Machine Learning Research (Accessed: 14 April 2026).
- [] Sciutti, A., Beetz, M., Inamura, T., et al. (2023) ‘The Present and the Future of Cognitive Robotics’, *IEEE Robotics & Automation Magazine*, 30(3), pp. 160–163. Available at: <https://ieeexplore-ieee.org.plymouth.idm.oclc.org/document/10255092> (Accessed: 18 March 2026).
- [] Smedegaard, C. V. (2019) ‘Reframing the Role of Novelty within Social HRI: From Noise to Information’, in *Proceedings of the 14th ACM/IEEE International Conference on Human-Robot Interaction (HRI ’19)*, pp. 411–420. Available at: <https://dl.acm.org/doi/10.1109/HRI.2019.8673219> (Accessed: 22 March 2026).
- ☒ [] Spezialetti, M., Placidi, G. and Rossi, S. (2020) ‘Emotion Recognition for Human-Robot Interaction: Recent Advances and Future Perspectives’, *Frontiers in Robotics and AI*, 7, Article 532279. Available

at: https://www.researchgate.net/publication/347520928_Emotion_Recognition_for_Human-Robot_Interaction_Recent_Advances_and_Future_Perspectives (Accessed: 14 April 2026).

- [] Tapus, A., Matarić, M. J. and Scassellati, B. (2007) ‘Socially assistive robotics [Grand Challenges of Robotics]’, *IEEE Robotics & Automation Magazine*, 14(1), pp. 35–42. Available at: <https://sczlab.yale.edu/sites/default/files/files/Tapus-RAM2007.pdf> (Accessed: 25 March 2026).

2.7. Appendix

2.7.1. Code

2.7.2. Video Demo

- YouTube link: test

Appendices

Appendix A: AI Declaration

Solo Work	S1 - Generative AI tools have not been used for this assessment.
Assisted Work	A1 – Idea Generation and Problem Exploration Used to generate project ideas, explore different approaches to solving a problem, or suggest features for software or systems. Students must critically assess AI-generated suggestions and ensure their own intellectual contributions are central.
	A2 - Planning & Structuring Projects AI may help outline the structure of reports, documentation and projects. The final structure and implementation must be the student's own work.
	A3 – Code Architecture AI tools maybe used to help outline code architecture (e.g. suggesting class hierarchies or module breakdowns). The final code structure must be the student's own work.
	A4 – Research Assistance Used to locate and summarise relevant articles, academic papers, technical documentation, or online resources (e.g. Stack Overflow, GitHub discussions). The interpretation and integration of research into the assignment remain the student's responsibility.
	A5 - Language Refinement Used to check grammar, refine language, improve sentence structure in documentation not code. AI should be used only to provide suggestions for improvement. Students must ensure that the documentation accurately reflects the code and is technically correct.
	A6 – Code Review AI tools can be used to check comments within the code and to suggest improvements to code readability, structure or syntax. AI should be used only to provide suggestions for improvement. Students must ensure that the code accurately reflects their knowledge and is technically correct.
	A7 - Code Generation for Learning Purposes Used to generate example code snippets to understand syntax, explore alternative implementations, or learn new programming paradigms. Students must not submit AI-generated code as their own and must be able to explain how it works.
	A8 - Technical Guidance & Debugging Support AI tools can be used to explain algorithms, programming concepts, or debugging strategies. Students may also help interpret error messages or suggest possible fixes. However, students must write, test, and debug their own code independently and understand all solutions submitted.
	A9 - Testing and Validation Support AI may assist in generating test cases, validating outputs, or suggesting edge cases for software testing. Students are responsible for designing comprehensive test plans and interpreting test results.
	A10 - Data Analysis and Visualization Guidance AI tools can help suggest ways to analyse datasets or visualize results (e.g. recommending chart types or statistical methods). Students must perform the analysis themselves and understand the implications of the results.
	A11 - Other uses not listed above Please specify:
Partnered Work	P1 - Generative AI tool usage has been used integrally for this assessment Students can adopt approaches that are compliant with instructions in the assessment brief. Please Specify:

Figure 5: Student Declaration of AI Tool use in this Assessment Table

I declare that I've used the AI tools listed below whilst preparing this assessment. I've read and understood the University of Plymouth's policy on the use of AI tools in assessment and confirm that my use falls within the coursework's allowed categories, i.e. **A2 (Planning and Structuring Projects)** and **A4 (Research Assistance)**.

AI Tool Used	Purpose of Use	Extent of Use
ChatGPT	Brainstorming project ideas and structuring the report (A2)	Initial brainstorming and final outline stages
ChatGPT	Reviewing structural alignment against grading criteria and mapping word-count budgets (A2)	After and midway through section-drafting
ChatGPT	Finding relevant pages to read in the paper (A4)	Few times if the paper is too long
ChatGPT	General conversations via web-search AI about prevalent papers to read about how the topics relates to others' studies (A4)	Few times at the end

- ☒ I understand that the ownership and responsibility for the academic integrity of this submitted assessment falls with me, the student.
- ☒ I confirm that all details provided above are an accurate description of how AI was used for this assessment.